assignment 07

Kiran Komati

5/16/2021

```
knitr::opts_chunk$set(echo = FALSE)
knitr::opts_knit$set(root.dir = 'C:/Users/kiran/dsc520')
```

Load the `data/r4ds/heights.csv` to heights_df

```
heights_df <- read.csv("data/r4ds/heights.csv")
```

Fit a linear model

```
earn_lm <-  lm(earn ~ ed + race + height + age + sex, data=heights_df)
```

View the summary of your model

```
summary(earn_lm)
```

```
##
## Call:
## lm(formula = earn ~ ed + race + height + age + sex, data = heights_df)
##
## Residuals:
##    Min     1Q Median     3Q    Max
## -39423  -9827  -2208   6157 158723
##
## Coefficients:
##               Estimate Std. Error t value Pr(>|t|)
## (Intercept)   -41478.4    12409.4  -3.342 0.000856 ***
## ed              2768.4      209.9  13.190  < 2e-16 ***
## racehispanic   -1414.3     2685.2  -0.527 0.598507
## raceother        371.0     3837.0   0.097 0.922983
## racewhite       2432.5     1723.9   1.411 0.158489
## height           202.5      185.6   1.091 0.275420
## age              178.3       32.2   5.537 3.78e-08 ***
## sexmale        10325.6     1424.5   7.249 7.57e-13 ***
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 17250 on 1184 degrees of freedom
## Multiple R-squared:  0.2199, Adjusted R-squared:  0.2153
## F-statistic: 47.68 on 7 and 1184 DF,  p-value: < 2.2e-16
```

Creating predictions using `predict()`

```r
predicted_df <- data.frame(
  earn = predict(earn_lm, newdata = heights_df),
  ed=heights_df$ed, race=heights_df$race, height=heights_df$height,
  age=heights_df$age, sex=heights_df$sex
  )
```

Compute deviation (i.e. residuals)

```r
mean_earn <- mean(heights_df$earn)
```

Corrected Sum of Squares Total

```r
sst <- sum((mean_earn - heights_df$earn)^2)
```

Corrected Sum of Squares for Model

```r
ssm <- sum((mean_earn - predicted_df$earn)^2)
```

Residuals

```r
residuals <- heights_df$earn - predicted_df$earn
```

Sum of Squares for Error

```r
sse <- sum(residuals^2)
```

R Squared

```r
r_squared <- ssm/sst
```

Number of observations

```
n <- nrow(heights_df)
```

Number of regression paramaters

```
p <- 8
```

Corrected Degrees of Freedom for Model

```
dfm <- p-1
```

Degrees of Freedom for Error

```
dfe <- n-p
```

Corrected Degrees of Freedom Total: DFT = n - 1

```
dft <- n-1
```

Mean of Squares for Model: MSM = SSM / DFM

```
msm <- ssm/dfm
```

Mean of Squares for Error: MSE = SSE / DFE

```
mse <- sse/dfe
```

Mean of Squares Total: MST = SST / DFT

```
mst <- sst/dft
```

F Statistic

```
f_score <- msm/mse
```

Adjusted R Squared R2 = 1 - (1 - R2)(n - 1) / (n - p)

```r
adjusted_r_squared <- 1 - (1 - r_squared)*(n - 1) / (n - p)
```